

Installing and using the Strimzi Operator with minikube

Install minikube

Make certain you have **kubect1** installed. You can install kubect1 according to the instructions in [Install and Set Up kubect1](#).

If you do not already have a hypervisor installed, install one of these now:

- [KVM](#), which also uses QEMU
- [VirtualBox](#)

Download the latest version of minikube

<https://github.com/kubernetes/minikube/releases>

```
$ sudo dnf install <path to>minikube-1.9.2-0.x86\_64.rpm
```

or

```
$ curl -LO minikube
```

```
https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64 \
&& chmod +x minikube
```

```
$ sudo mkdir -p /usr/local/bin/
```

```
$ sudo install minikube /usr/local/bin/
```

```
$ minikube version
```

```
minikube version: v1.9.2
```

```
commit: 93af9c1e43cab9618e301bc9fa720c63d5efa393
```

See <https://kubernetes.io/docs/tasks/tools/install-minikube/> for additional info.

Start minikube

You can choose the driver, options are virtualbox, kvm, and podman (experimental).

```
$ minikube start --driver=virtualbox --kubernetes-version=1.18.0
```

```
😊 minikube v1.9.2 on Fedora 31
```

```
✨ Using the virtualbox driver based on user configuration
```

```
📀 Downloading VM boot image ...
```

```
> minikube-v1.9.0.iso.sha256: 65 B / 65 B [-----] 100.00% ? p/s  
0s
```

```
> minikube-v1.9.0.iso: 174.93 MiB / 174.93 MiB [-] 100.00% 9.20 MiB p/s  
20s
```

```
👍 Starting control plane node m01 in cluster minikube
```

```
📁 Downloading Kubernetes v1.18.0 preload ...
```

```
> preloaded-images-k8s-v2-v1.18.0-docker-overlay2-amd64.tar.lz4: 542.91  
MiB
```

```
🔥 Creating virtualbox VM (CPUs=2, Memory=5900MB, Disk=20000MB) ...
```

```
🐳 Preparing Kubernetes v1.18.0 on Docker 19.03.8 ...
```

```
🌟 Enabling addons: default-storageclass, storage-provisioner
```

```
🏡 Done! kubectl is now configured to use "minikube"
```

```
$ minikube status
```

```
m01
```

```
host: Running
```

```
kubelet: Running
```

```
apiserver: Running
```

```
kubeconfig: Configured
```

Install the Strimzi Operator

```
$ curl -LO https://strimzi.io/install/latest
```

```
$ curl -LO https://strimzi.io/install/latest | sed 's/namespace: ./namespace:  
default/' | kubectl apply -f -
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
strimzi-cluster-operator-6c8d574d49-hch89	1/1	Running	0	91s

Create the Kafka Cluster

```
$ curl -LO https://strimzi.io/examples/latest/kafka/kafka-persistent.yaml | vi -
```

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    version: 2.4.0
    replicas: 3
    listeners:
      plain: {}
      tls: {}
      external:
        type: loadbalancer
        tls: false
    config:
      offsets.topic.replication.factor: 3
      transaction.state.log.replication.factor: 3
      transaction.state.log.min.isr: 2
      log.message.format.version: "2.4"
    storage:
      type: jbod
      volumes:
        - id: 0
          type: persistent-claim
          size: 100Gi
          deleteClaim: false
  zookeeper:
    replicas: 3
    storage:
      type: persistent-claim
      size: 100Gi
      deleteClaim: false
  entityOperator:
    topicOperator: {}
    userOperator: {}
:w ! kubectl apply -f -

kafka.kafka.strimzi.io/my-cluster created
```

Use 'minikube tunnel' to expose external IP

`minikube tunnel` runs as a process, creating a network route on the host to the service CIDR of the cluster using the cluster's IP address as a gateway. The tunnel command exposes the external IP directly to any program running on the host operating system.

In another terminal, run:

```
$ minikube tunnel
```

Status:

```
  machine: minikube
  pid: 112455
  route: 10.96.0.0/12 -> 192.168.99.102
  minikube: Running
  services: [my-cluster-kafka-0, my-cluster-kafka-1, my-cluster-kafka-2,
my-cluster-kafka-external-bootstrap]
  errors:
    minikube: no errors
    router: no errors
    loadbalancer emulator: no errors
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
my-cluster-entity-operator-5dd6d7f5bd-c76b4	3/3	Running	0	88m
my-cluster-kafka-0	2/2	Running	0	88m
my-cluster-kafka-1	2/2	Running	0	88m
my-cluster-kafka-2	2/2	Running	0	88m
my-cluster-zookeeper-0	2/2	Running	0	90m
my-cluster-zookeeper-1	2/2	Running	0	90m
my-cluster-zookeeper-2	2/2	Running	0	90m
strimzi-cluster-operator-6c8d574d49-hch89	1/1	Running	0	107m

```
$ kubectl get kafka -o yaml
```

```
apiVersion: v1
items:
- apiVersion: kafka.strimzi.io/v1beta1
  kind: Kafka
  metadata:
    annotations:
      kubectl.kubernetes.io/last-applied-configuration: |

{"apiVersion":"kafka.strimzi.io/v1beta1","kind":"Kafka","metadata":{"annotations":{},"name":"my-cluster","namespace":"default"},"spec":{"entityOperator":{"topicOperator":{},"userOperator":{}},"kafka":{"config":{"log.message.format.version":"2.4","offsets.topic.replication.factor":3,"transaction.state.log.min.isr":2,"transaction.state.log.replication.factor":3},"listeners":{"external":{"tls":false,"type":"loadbalancer"},"plain":{},"tls":{}},"replicas":3,"storage":{"type":"jbod","volumes":[{"deleteClaim":false,"id":0,"size":"100Gi","type":"persistent-claim"}]},"version":"2.4.0"},"zookeeper":{"replicas":3,"storage":{"deleteClaim":false,"size":"100Gi","type":"persistent-claim"}}}}
    creationTimestamp: "2020-04-15T14:40:27Z"
    generation: 1
    managedFields:
    - apiVersion: kafka.strimzi.io/v1beta1
      fieldsType: FieldsV1
      fieldsV1:
        f:spec:
          f:entityOperator:
            f:topicOperator:
              f:reconciliationIntervalSeconds: {}
              f:topicMetadataMaxAttempts: {}
              f:zookeeperSessionTimeoutSeconds: {}
            f:userOperator:
              f:reconciliationIntervalSeconds: {}
              f:zookeeperSessionTimeoutSeconds: {}
          f:status:
            f:conditions: {}
            f:listeners: {}
            f:observedGeneration: {}
        manager: okhttp
        operation: Update
        time: "2020-04-15T14:43:17Z"
      name: my-cluster
      namespace: default
      resourceVersion: "4802"
      selfLink:
      /apis/kafka.strimzi.io/v1beta1/namespaces/default/kafkas/my-cluster
      uid: 55993681-f909-40ad-ba4c-b94968aba11f
    spec:
```

```
entityOperator:
  topicOperator: {}
  userOperator: {}
kafka:
  config:
    log.message.format.version: "2.4"
    offsets.topic.replication.factor: 3
    transaction.state.log.min.isr: 2
    transaction.state.log.replication.factor: 3
  listeners:
    external:
      tls: false
      type: loadbalancer
    plain: {}
    tls: {}
  replicas: 3
  storage:
    type: jbod
    volumes:
      - deleteClaim: false
        id: 0
        size: 100Gi
        type: persistent-claim
  version: 2.4.0
zookeeper:
  replicas: 3
  storage:
    deleteClaim: false
    size: 100Gi
    type: persistent-claim
status:
  conditions:
    - lastTransitionTime: 2020-04-15T14:43:17+0000
      status: "True"
      type: Ready
  listeners:
    - addresses:
        - host: my-cluster-kafka-bootstrap.default.svc
          port: 9092
      type: plain
    - addresses:
        - host: my-cluster-kafka-bootstrap.default.svc
          port: 9093
  certificates:
    - |
      -----BEGIN CERTIFICATE-----
      MIIDLTCCAhWgAwIBAg...
      ...
```

```
-----END CERTIFICATE-----
  type: tls
- addresses:
  - host: 10.102.68.38    //Note the IP address
    port: 9094           //Note the port
    type: external
  observedGeneration: 1
kind: List
metadata:
  resourceVersion: ""
  selfLink: ""
```

Create a Kafka topic

Change the partitions and replicas to **3** for a more compelling demo.

```
$ curl -LO https://strimzi.io/examples/latest/topic/kafka-topic.yaml | vi -
```

```
apiVersion: kafka.strimzi.io/v1beta1
```

```
kind: KafkaTopic
```

```
metadata:
```

```
  name: my-topic
```

```
  labels:
```

```
    strimzi.io/cluster: my-cluster
```

```
spec:
```

```
  partitions: 3
```

```
  replicas: 3
```

```
  config:
```

```
    retention.ms: 7200000
```

```
    segment.bytes: 1073741824
```

```
:w ! kubectl apply -f -
```


Edit and run the Producer/Consumer (see files below)

```
$ cd <path to source>
```


```
$ code .
```

Add the server IP and port from the output of the 'kubectl get kafka -o yaml' command to both Producer.java (line 31) and Consumer.java (line 32):

```
-----END CERTIFICATE-----
  type: tls
- addresses:
  - host: 10.102.68.38
    port: 9094
  type: external
observedGeneration: 1
kind: List
metadata:
  resourceVersion: ""
  selfLink: ""

/*
 * Producer configuration
 */
    Map<String, Object> props = new HashMap();
    props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "10.102.68.38:9094");

/*
 * Consumer configuration
 */
    Map<String, Object> props = new HashMap();
    props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "10.102.68.38:9094");
```



Run `Producer.java`, notice the message was sent to partition 1, offset 1:

```
[main] INFO io.strimzi.demo.Producer - Message sent (topic my-topic, partition 1, offset 1)
```

Run `Consumer.java`, notice the 'Hello World' message was received:

```
[main] INFO io.strimzi.demo.Consumer - Received message: null / Hello World (from topic my-topic, partition 1, offset 1)
```

Stop minikube

```
$ minikube stop
```



```
Stopping "minikube" in virtualbox ...
```



```
Node "m01" stopped.
```

View the minikube dashboard

\$ minikube dashboard

😞 Verifying dashboard health ...

🚀 Launching proxy ...

😞 Verifying proxy health ...

🎉 Opening

http://127.0.0.1:36805/api/v1/namespaces/kubernetes-dashboard/services/http:kubernetes-dashboard:/proxy/ in your default browser...

[205204:205204:0416/155240.945361:ERROR:edid_parser.cc(102)] Too short EDID data: manufacturer id

[205204:205204:0416/155240.945618:ERROR:edid_parser.cc(102)] Too short EDID data: manufacturer id

Opening in existing browser session.

The screenshot shows the Kubernetes Dashboard Overview page. The left sidebar contains navigation links for Cluster, Namespaces, Nodes, Persistent Volumes, Storage Classes, Workloads, Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets, and Discovery and Load Balancing. The main content area is titled 'Workloads' and 'Workload Status'. It features four large green circles representing the status of Deployments, Pods, Replica Sets, and Stateful Sets. Below this, there is a table titled 'Deployments' showing a list of deployments in the 'default' namespace. The table columns are Name, Namespace, Labels, Pods, Age, and Images. Two deployments are listed: 'my-cluster-entity-operator' and 'strimzi-cluster-operator', both with 1/1 pods and an age of 'a day'. The 'strimzi-cluster-operator' deployment is highlighted with a green checkmark. The bottom of the page shows pagination information: '1 ~ 2 of 2'.

Name	Namespace	Labels	Pods	Age	Images
my-cluster-entity-operator	default	app.kubernetes.io/instance: my-cluster	1 / 1	a day	strimzi/operator:0.17.0
strimzi-cluster-operator	default	app.kubernetes.io/managed-by: strimzi-cluster-operator	1 / 1	a day	strimzi/operator:0.17.0